

Image compression using neural auto-encoder and quantization

License MIT (<https://opensource.org/licenses/MIT>)

This project is a simple implementation of auto-encoder neural network for image compression. The auto-encoder neural network is trained on the ImageNet dataset. The trained model is then used to compress and decompress the images.

Navigation:

- [Model architecture](#)
- [Download pretrained models](#)
- [Quick start](#)
- [Compression](#)
- [Decompression](#)
- [Training from scratch](#)
- [Results](#)
- [Notebooks](#)

Model architecture

Model represents a variational auto-encoder with residual blocks and skip connections.

- Encoder: *ResNet-18 architecture with fully connected layers*
- Decoder: *ResNet-18 architecture with transposed convolution layers*
- Loss: *VGG loss + MSE loss*

Download pretrained models

Models were trained on [130k Images \(512x512\) - Universal Image Embeddings](https://www.kaggle.com/datasets/rhtsingh/130k-images-512x512-universal-image-embeddings) (<https://www.kaggle.com/datasets/rhtsingh/130k-images-512x512-universal-image-embeddings>) dataset from Kaggle.

Here are the links to download the pretrained models:

B = number of quantization levels

- [B=2, resnet18](https://drive.google.com/drive/folders/1FaeWzeRW3BMqqZwGsHUjhf7PuAOsiY6E?usp=sharing) (<https://drive.google.com/drive/folders/1FaeWzeRW3BMqqZwGsHUjhf7PuAOsiY6E?usp=sharing>)
- [B=8, resnet18](https://drive.google.com/drive/folders/1fYDc0e43cUR7xslYatpz8fdJ_6KMJmSs?usp=sharing) (https://drive.google.com/drive/folders/1fYDc0e43cUR7xslYatpz8fdJ_6KMJmSs?usp=sharing)

Put downloaded models in models directory.

Quick start

[compress_all.sh \(scripts/compress_all.sh\)](#) compresses all images from assets/images directory and saves them in assets/compressed directory.

compress_all.sh takes 3 arguments:

- qb - number of quantization levels
- resnet-model - resnet model architecture
- device - torch device to evaluate on

```
# Compress all images from assets/images directory
bash scripts/compress_all.sh 8 resnet18 cpu
```

[decompress_all.sh \(./scripts/decompress_all.sh\)](#) decompresses all images from assets/compressed directory and saves them in assets/decompressed directory.

decompress_all.sh takes 3 arguments:

- qb - number of quantization levels
- resnet-model - resnet model architecture
- device - torch device to evaluate on

```
# Decompress all images from assets/compressed directory
bash scripts/decompress_all.sh 8 resnet18 cpu
```

Compression

```
# Compress the `baboon` image from assets/images directory
python compress.py \
  --image=assets/images/baboon.png \
  --output=assets/compressed/baboon.bin \
  --models-dir=models \
  --resnet-model=resnet18 \
  --qb=8 \
  --device=cuda
```

Decompression

```
# Decompress the compressed image
python decompress.py \
  --file=assets/compressed/baboon.bin \
  --output=assets/decompressed/baboon.png \
  --qb=8 \
  --resnet-model=resnet18 \
  --models-dir=models \
  --device=cuda
```

Training from scratch


```
python train.py \
--root [path to images] \
--test-root [path to test images] \
--resnet-model [resnet model architecture] \
--qb [number of quantization levels] \
--epochs [number of epochs] \
--batch-size [batch size] \
--lr [learning rate] \
--device [torch device to train on] \
--save-results-every [save results every n epochs] \
--save-models-dir [path to save models] \
--use-checkpoint [use checkpoint to resume training]
```

Results

Size comparison

Image	Original (.png)	Jpeg	B=2	B=8
baboon (assets/images/baboon.png)	624 KB	100 KB	20 KB	76 KB
lena (assets/images/lena.png)	504 KB	62 KB	20 KB	76 KB
peppers (assets/images/peppers.png)	528 KB	56 KB	20 KB	76 KB


Images






Graphs

PSNR: Peak signal-to-noise ratio



BPP: Bits per pixel



Quality comparison:

Original





JPG



B=8



B=2



Notebooks

- [Kaggle training notebook \(notebooks/kaggle-cuda-training.ipynb\)](#)
- [Analysis notebook \(notebooks/analysis.ipynb\)](#)